

Sriram Balasubramanian

📍 College Park, MD ✉ sriramb@umd.edu 🌐 www.sriram.live in sriramb1998 📄 Google Scholar

Education

University of Maryland, College Park <i>PhD in Computer Science (advised by Prof. Soheil Feizi 🔗)</i>	August 2023 – December 2026 (Expected) GPA: 4.0
University of Maryland, College Park <i>MS in Computer Science (advised by Prof. Soheil Feizi 🔗)</i>	August 2021 – May 2023 GPA: 4.0
Indian Institute of Technology, Bombay <i>B. Tech (Hons) in Computer Science & Engineering (advised by Prof. Sunita Sarawagi 🔗)</i>	August 2016 – May 2020 GPA: 9.56/10.0

Work Experience

Research Fellow (Incoming) <i>FAR.AI</i>	Berkeley, CA May 2026 – August 2026
Research Intern <i>Adobe (Document Intelligence Lab)</i>	San Jose, CA May 2025 – August 2025
- Created a recipe to curate data and <i>post-train</i> LLMs using SFT and RL based algorithms to generate <i>citations</i> to answers by utilizing a decomposition based strategy. Preprint available here 🔗. - Designed a new prompting strategy and a prompt optimization pipeline to improve attribution methods in Acrobat AI assistant by up to 10%.	
AI Researcher <i>RelAI 🔗</i>	College Park, MD June 2023 – August 2023 June 2024 – August 2024
- Developed and designed multiple applications to enhance vision model reliability for RelAI - Involved in the development of RelAI, contributing from the inception stage through planning, execution, and deployment phases.	
Research Intern <i>Comcast</i>	Washington D.C. June 2022 – August 2022
- Investigated the effectiveness of transfer learning in deep neural networks in the low resource regime (when the target domain has very limited data). - Devised non-neural methods which could outperform both traditional collaborative filtering methods and neural networks in this regime.	
Research Fellow <i>Microsoft Research</i>	Bangalore, India August 2020 – August 2021
Predicting e-mail arrivals and reads: Built machine learning models to predict e-mail arrivals and reads from user type and history of arrivals/reads to improve cache hit rates. Simulating network paths using ML: Built machine learning models to simulate internet paths using static network traces	
Intern <i>Tower Research Capital</i>	Gurgaon, India May 2019 – August 2019
Developed tools to analyze trading and market logs, and compute and display metrics for various trading strategies and securities.	

Publications and Preprints

Retrieval is Enough: Training-Free Interpretability with a Tool-Using Agent S Balasubramanian, S Feizi <i>Under review at NeurIPS, (link here) 🔗, 2026</i>
Retrieval is Enough: Training-Free Interpretability with a Tool-Using Agent HK Tan*, S Kabir*, S Agrawal, S Chereddy, S Balasubramanian

Mentoring at SPAR, In Progress ([link here](#)) [🔗](#), 2026

Decomposition-Enhanced Training for Post-Hoc Attributions In Language Models

S Balasubramanian, S Basu, K Goswami, R Rossi, V Manjunatha, R Santhosh, R Zhang, S Feizi, N Lipka
European Chapter of the Association for Computational Linguistics, Main, 2026

A Closer Look at Bias and Chain-of-Thought Faithfulness of Large (Vision) Language Models

S Balasubramanian, S Basu, S Feizi

Empirical Methods in Natural Language Processing (EMNLP), Findings, 2025

Decomposing and Interpreting Image Representations via Text in ViTs Beyond CLIP

S Balasubramanian, S Basu, S Feizi

Spotlight at *Mechanistic Interpretability Workshop, ICML*, 2024

Advances in Neural Information Processing Systems (NeurIPS), 2024

Exploring Geometry of Blind Spots in Vision Models

S Balasubramanian*, G Sriramanan*, VS Sadasivan, S Feizi

Spotlight at *Advances in Neural Information Processing Systems (NeurIPS)*, 2023

Towards Improved Input Masking for Convolutional Neural Networks

S Balasubramanian, S Feizi

IEEE/CVF International Conference on Computer Vision (ICCV), 2023

What's in a Name? Are BERT Named Entity Representations just as Good for any other Name?

S Balasubramanian*, N Jain*, G Jindal*, A Awasthi, S Sarawagi

Rep4NLP Workshop @ Annual Meeting of the Association of Computational Linguistics (ACL), 2020

Can AI-Generated Text be Reliably Detected? Stress Testing AI Text Detectors Under Various Attacks

VS Sadasivan, A Kumar, **S Balasubramanian**, W Wang, S Feizi

Transactions on Machine Learning Research (TMLR), 2025

Media coverage at [Washington Post](#) [🔗](#), [Wired](#) [🔗](#), [TechSpot](#) [🔗](#), [New Scientist](#) [🔗](#)

Rethinking Copyright Infringements in the Era of Text-to-Image Generative Models

M Moayeri, **S Balasubramanian**, S Basu, P Kattakinda, A Chegini, R Brauneis, S Feizi

International Conference on Learning Representations (ICLR), 2025

Gaming Tool Preferences in Agentic LLMs

K Faghieh, W Wang, Y Cheng, S Bharti, G Sriramanan, **S Balasubramanian**, P Hosseini, S Feizi

Empirical Methods in Natural Language Processing (EMNLP), Main Conference, 2025

Simulating Network Paths with Recurrent Buffering Units

D Anshumaan*, **S Balasubramanian***, S Tiwari, N Natarajan, S Sellamanickam, VN Padmanabhan

AAAI Conference on Artificial Intelligence (AAAI), 2023

Hop, Skip, and Overthink: Diagnosing Why Reasoning Models Fumble during Multi-Hop Analysis

A Yadav, I Nalawade, S Pillarichety, Y Babu, R Ghosh, S Basu, W Zhao, W Zhao, A Nasaeh, **S Balasubramanian**, S Srinivasan

arXiv preprint arXiv:2508.04699, 2025

Seeing What's Not There: Spurious Correlation in Multimodal LLMs

P Hosseini, S Nawathe, M Moayeri, **S Balasubramanian**, S Feizi

arXiv preprint arXiv:2503.08884, 2025

A Survey on Mechanistic Interpretability for Multi-Modal Foundation Models

Z Lin, S Basu, M Beigi, V Manjunatha, RA Rossi, Z Wang, Y Zhou, Y Zhou, **S Balasubramanian**, A Zarei, K Rezaei, Y Shen, B Menglong Yao, Z Xu, Q Liu, Y Zhang, Y Sun, S Liu, L Shen, H Li, S Feizi, L Huang

arXiv preprint arXiv:2502.17516, 2025

Services and Teaching

- **SPAR [↗](#) Mentor**: Mentoring 4 researchers (across multiple timezones) on projects related to utilizing circuits for solving AI safety problems
- Reviewer for prominent machine learning conferences such as ICML 2024, NeurIPS 2024 (Top Reviewer), ICLR 2025, NeurIPS 2025 (Top Reviewer), ICLR 2026, NeurIPS 2026.
- Introduced high-school students to AI as an instructor as part of **TRAILS AI Summer Camp [↗](#)**
- Teaching Assistant for Programming Handheld Systems (CMSC 436), Probability and Statistics (STAT 400) at UMD College Park; and Data Interpretation and Analysis (CS 215) and Electricity and Magnetism (PH 108) at IIT Bombay.

Awards and Honors

- Awarded Dean's Fellowship at the University of Maryland for the first year of PhD [2023]
- Awarded Institute Academic Prize for exceptional academic performance in IIT Bombay [2017]
- Ranked **2nd** in the institute out of about 900 students in the first year at IIT Bombay [2017]
- Ranked **4th** in JEE Mains out of 1.2 million candidates all over India [2016]
- Ranked **2nd** in the Maharashtra State Board Examinations (12th grade) [2016]